

[Open in app ↗](#)

° + ❀ A Cozy Walk Through the History of AI ❤️



Divya karlapudi 10 min read · Nov 22, 2025



...

A journey from neurons to transformers ❀

★ ° ⓘ Introduction:

When we think of AI today, we picture ChatGPT answering questions like a friend, self-driving cars navigating traffic, and deepfake videos that look scarily real. But AI didn't just appear one morning. It wasn't magic. It was built, one idea at a time.

For decades, researchers often working quietly in labs kept pushing the boundaries of imagination and computation. And yet, when we try to learn AI today, it's easy to feel overwhelmed. There are too many buzzwords, too many frameworks, too many "must-learn" topics.

So... what if we start from the beginning instead? ^ ^ ^

↳ Where did these ideas come from?

↳ Why were they created?

↳ What problems were researchers trying to solve?

Understanding the journey, not just the destination makes AI surprisingly human, relatable, and even comforting. (yes, comforting (◦•◦)) ❤️

💡 If we first understand the problems that existed.. the solutions begin to make sense.

So, let's start with history. The ideas, the inspirations, the breakthroughs.

Let's understand how AI slowly went from theory... to code... to reality 😊

Ready? 🤓

⚡ It was 1943.. the first artificial neuron was born.

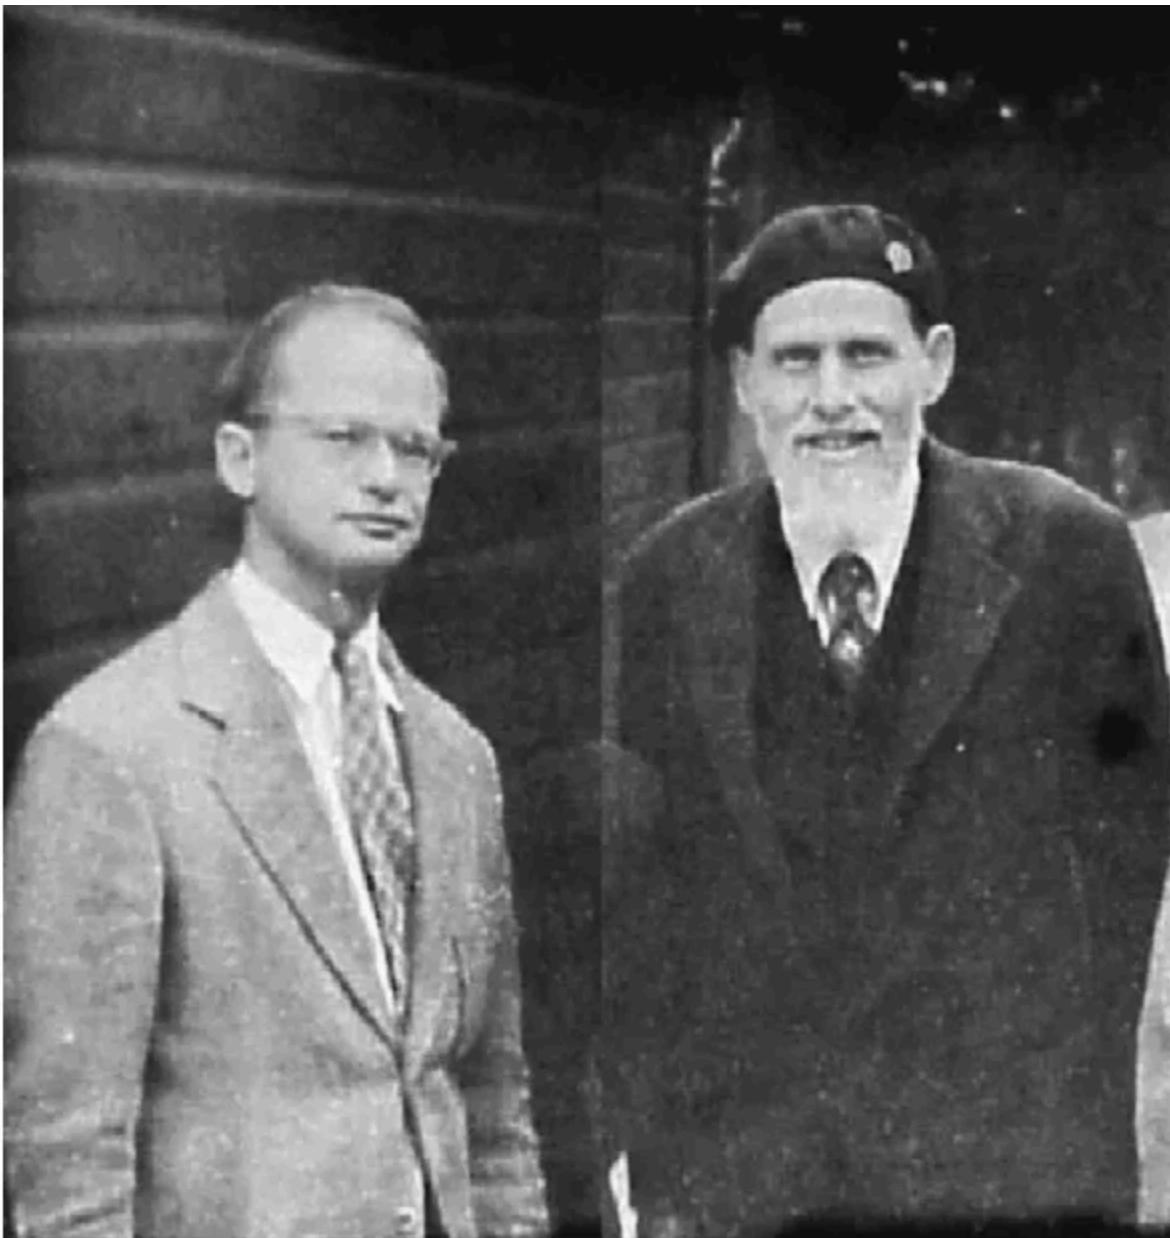
Sounds unbelievable, right?

1943. No internet. No GPUs. No Python.

Yet... the first concept of an artifical neuron was already on paper 🌟

That year, Warren McCulloch (a neurophysiologist) and Walter Pitts (a young mathematician with no formal degree) wrote a groundbreaking paper titled:

“A Logical Calculus of the Ideas Immanent in Nervous Activity”



Walter Pitts (left) & McCulloch (Right)

This was the birth of the first mathematical model of a neuron (**McCulloch-Pitts neuron**).

It didn't learn. It didn't adapt. But it was logical (literally). It used AND, OR, NOT operations just like today's digital circuits.

Why was this a big deal (😲)?

Because it showed something revolutionary:

(💡 Maybe the brain could be explained using logic.
And if the brain is logical... maybe a machine could think too.)

This was the first time anyone treated **thinking** as something mathematically computable.

👉 *This paper lit the first spark of Artificial Intelligence.*

Here's the actual paper '🕒 : A Logical Calculus of the Ideas Immanent in Nervous Activity (1943)'

From this humble beginning... the road to AI officially started to unfold ✨

—99—

✨ 1958. The perceptron brings hope ☺

It's the late 1950s. Computers are huge, expensive, and painfully slow.

AI isn't even a defined term yet.

But then, a young psychologist at Cornell named Frank Rosenblatt introduces something bold... (He was originally a psychologist, not a computer scientist >v<)

A machine that could learn from data. A machine inspired by how the human brain works.

A machine called “The perceptron” 🎉

💡 Why Was This a Big Deal?

For the first time in history, a system could learn patterns from inputs and make decisions.

🌸 The perceptron was basically an early **single-layer neural network** and, the idea was beautiful:

It could **recognize simple shapes, classify patterns, and adjust its weights** over time. Which means... **It could learn.** (That word “learn” changed everything.)

What Could the Perceptron Do? 🤔

Let's say we show it multiple samples of **circles and triangles**.

Over time, it learns to say: “**Yes, this looks like a circle**” Or, “**Nope.. that's more like a triangle.**”

Sounds basic today. But in 1958, this felt like magic 😲

This was the first step towards modern neural networks.

Frank Rosenblatt

The perceptron was powerful, but there was a problem 🤔 :

It couldn't learn complex patterns like XOR functions.

Here's the original paper ↗ : “The Perceptron: A Probabilistic Model for Information Storage and Organization in the Brain”

Right now, in 1958, AI was full of hope. Our, Next Stop 1960s & 1970s: The First AI Winter ❄️

—————9—————

Let's slow down the timeline and feel what happened in 1969, because this moment didn't just slow AI down... It froze it ❄️

❄️ 1969. The First AI Winter Begins...

The excitement from Rosenblatt's perceptron was still echoing across research labs.

Many believed machines would soon think, learn, even see like humans.

But then...

Two researchers “Marvin Minsky” and “Seymour Papert” published a book that would reshape AI history.

■ “Perceptrons” (1969): A mathematical and critical analysis of neural networks.

Minsky & Papert mathematically proved something devastating:

💡 Single-layer perceptrons cannot learn certain simple functions (like XOR).

That means:

↳ They could only separate things **linearly**.

↳ If data was **complex or overlapping?** → It failed

↳ It couldn't understand relations like:

“If both inputs are different, output should be 1.” (XOR)

In short: The perceptron was more limited than everyone thought (•◦_◦)

Consequences? 🍀

Government funding agencies lost hope. Investors pulled back. Universities stopped encouraging AI research.

This period is known as the **First AI Winter (1969–1980s)** because, **interest froze. Funding froze. Careers froze.**

Many researchers moved on to other fields like **Statistics, Expert systems...**
Neural networks were abandoned 😞

Marvin Minsky (left) and Seymour Papert (right)

Their critique was correct. But, only for single-layer neural networks. They didn't consider multi-layer networks (MLPs) Or, backpropagation because, these concepts weren't even discovered yet!

So.. AI wasn't dead. It was simply waiting... for the next breakthrough ❤️

Here's the original scan of the book '₹ : Perceptrons – An Introduction to Computational Geometry

—99—

↗ 1980s. AI winter is about to melt... ε(◦•◦) ↪ ❤️

After the criticism of perceptrons in 1969, many believed neural networks were useless.

But in Japan, one researcher wasn't ready to give up ✨

Kunihiko Fukushima. The Quiet Revolutionary (personal fav) .(> ^ <), ❤️

Kunihiko Fukushima

In 1980, Kunihiko Fukushima introduced something groundbreaking.

💡 **The Neocognitron.** (A multi-layered neural architecture deeply inspired by the human visual system).

This architecture became the ancestor of modern **Convolutional Neural Networks (CNNs)** ❤️

What Was the Neocognitron? 🤔

It was designed to recognize patterns like shapes and handwritten characters even if they were shifted, distorted, or scaled (well.. to a large extent).

It consisted of layers (where S-cells (simple cells), detect basic features like edges.. C-cells (complex cells), become invariant to position and shape.. Hierarchical structure.. Increasing abstraction across layers)

(No.. no.. don't get overwhelmed, you will get it if you read the paper)

👉 This structure is very similar to modern CNNs like LeNet-5 & ResNet!



Fig explaining the interconnections between layers. (from Neocognitron paper)

Here's the paper ↗ : “[Neocognitron: A Self-organizing Neural Network Model for a Mechanism of Pattern Recognition Unaffected by Shift in Position](#)”

Even though the Neocognitron was brilliant, It learned via local rules but wasn't fully efficient or scalable.

We were adjusting weights with no systematic method to minimize error.

👉 Imagine: We feed an image of a “cat” into our neural network, if it outputs “dog” i.e., total error is high but, the question is **Which weight caused this mistake? Which layer is at fault? How much should each weight change?**

Without an answer to these questions, learning was impossible. Especially in deep layers.

We needed a method that could **compute the error, trace it back through all the layers and adjust each weight precisely.**

Let's step into **1986**, a year that revived AI. It was like the moment someone found the missing key... and opened the door again 🚪 ✨

—99—

🚀 1986. Backpropagation Brings Neural Networks Back to Life

👨‍🔬 Three researchers “Geoffrey Hinton”, “David Rumelhart”, and “Ronald Williams” have officially introduced **backpropagation** as a practical training method for multi-layer neural networks.

How exactly backpropagation works?

1. Feed input forward (get prediction)
2. Compare with the correct answer (calculate error)
3. Send the error backwards through layers
4. Adjust each weight using calculus

5. Repeat until the error is small.

This allowed networks to actually learn, rather than guess.

What did they show? ⚡

They demonstrated that Multi-layer neural networks can learn complex patterns, Deep layers can be trained using gradient descent, Errors can be propagated backward to adjust every weight 🏆

‘This paper is considered a turning point in AI history’

Here’s the paper ↗ : “Learning representations by back-propagating errors”

————— 99 ———

“But, can a network remember? Can it handle sequences, language, time?” 😰

⭐ “What if neural networks could not only think... but also remember?”

Backpropagation (1986) changed everything. It allowed neural networks to finally learn, layer by layer.

But soon, researchers hit a new limitation: Standard neural nets looked at data as if every input was independent.

↳ A word has no connection to the previous word.

↳ A video frame knew nothing about the last frame.

↳ A time series point didn't know it's past.

That's not how we humans think.. We understand sequences, context, stories, time.

So researchers asked: “**Can we build a neural network that learns over time?**”

Recurrent Neural Networks (RNNs). First Step Towards Intelligence.

Between 1986 and 1990, Recurrent Neural Networks (RNNs) were introduced. Two major names appear here.

“**Michael I. Jordan**” : 1986 : Introduced the **Jordan Network** (a neural network that loops its output back as input)

Michael I. Jordan

“Jeffrey L. Elman” : 1990 : Introduced the Elman network (used a hidden “context layer” to preserve memory.

Jeffrey L. Elman

They both had the same core philosophy: “**Let the network remember its previous state**”

This wasn’t just brilliant. It was the first computational model of memory in AI 🎉

But, how RNNs work? 😊

Imagine reading a sentence. Word by word. Our brain doesn’t forget the first word when reading the second one.

RNNs try to replicate this. They take input, update an internal state, and pass it forward.

Input → Hidden State (memory) → Output → Next Hidden State → ...

🎉 This little loop became groundbreaking as it allowed AI to process Speech, Text, Time-series data, Sensor signals, Music...

🚧 But There Was a Problem... (you saw it coming right? 😅)

Training RNNs with regular backpropagation failed badly.

The gradients became:

1. Too small → network forgot everything (vanishing gradient)
2. Too large → network exploded (exploding gradient)

The Result? : RNNs could remember.. but only short-term memory.

When they try to remember information across many time steps, they fail.

Let's imagine this sentence:

“The book that the professor who I met yesterday recommended was fascinating.”

To understand “*was fascinating*,”
you need to remember “**the book...**” from *way earlier* in the sentence.

But, RNNs *forget too quickly*. As the sequence grows:

Old information fades. Only recent inputs are remembered.

Why? Because of vanishing gradients.

Here are the paper links :)

Michael I. Jordan (1986): “SERIAL ORDER: A PARALLEL DISTRIBUTED PROCESSING APPROACH”

Jeffrey L. Elman (1990): “Finding Structure in Time”

—99—

Welcome to the 1990s, the first time AI finally proved itself in the real world
★*°。*

🧠 1997: The Birth of Long Short-Term Memory (LSTMs)

By the mid 1990s, researchers had hit a wall. RNNs failed to handle long-term dependencies because of the vanishing gradient problem.

That's when "Sepp Hochreiter" & "Jurgen Schmidhuber" proposed a breakthrough architecture in 1997.

Long Short-Term Memory (LSTM): A neural network that learns what to remember... and what to forget. (How cool is that? (〃^▽^〃))

This changed AI forever. It enabled deep learning in language, speech, translation, music, handwriting, and later, paved the way for GPT and Transformers.

🔍 What Was the Core Idea?

Instead of blindly passing memory from step to step (like RNNs), LSTMs control memory using gates.

Just like how our brain filters important events and forgets irrelevant noise.

1. **Memory Cell:** Long-term memory.
2. **Gates:** Attention / decision making.
3. **Forget Gate:** "This isn't important."
4. **Input Gate:** "Store this!"
5. **Output Gate:** "Use this now."

LSTMs overcame two major limitations in RNNs. Vanishing gradients and rapid memory loss.

They introduced a new concept: **selective memory**, choosing what to forget and what to store.

Enabled neural networks to **retain information across dozens, even hundreds of time steps.** 🎉

Here's the paper ↗ : “**LONG SHORT-TERM MEMORY**”

—99—

The 1986 backpropagation breakthrough revived neural networks. But, scientists still asked: “Can it actually solve real problems?” 🍃

In the early 1990s, someone answered that question.

★ Yann LeCun. The Architect of Modern Computer Vision.

A small, circular portrait of Yann LeCun, a man with short dark hair and a slight smile, wearing a light-colored shirt.

Yann LeCun (Photo by Nick Fetter/The New York Academy of Sciences)

A passionate researcher at AT&T Bell Labs, he introduced a neural network that could read handwritten digits automatically.

This was LeNet-5. One of the first convolutional neural networks (CNNs) used in production.

What was so special about LeNet-5? 

It could read ZIP codes from envelopes. Completely automatically. Used by the U.S. Postal Service, no human intervention needed!

↳ **Convolutional Layers:** Detect patterns like edges & shapes.

↳ **Pooling Layers:** Make model resistant to shifts/distortions.

↳ **Backpropagation:** Enabled training end-to-end

↳ **Real-world data:** First practical AI success

This was a huge break-through °❀★.○◐*:·

Here's the paper : “[Gradient-Based Learning Applied to Document Recognition](#)” (1998)

Also, check out: “[Convolutional Networks and Applications in Vision](#)” (2023)

From ImageNet and AlexNet to GANs, ResNet, Transformers, Retrieval-Augmented Generation, and eventually LLM agents... AI evolved rapidly, step by step.

These exciting advancements deserve their own deep dives... and that's exactly what we'll explore in our upcoming articles ❤

I really hope this article gave you a comfortable and confident start on your AI journey ★ °✿+°

With Love 🌟~..♡

Divya Karlapudi.

History Of Ai

Artificial Intelligence

Rnn

Cnn

Lstm Model



Written by **Divya karlapudi**

3 followers · 33 following

Edit profile

I love sharing what's in my mind in general.

No responses yet



...



Divya karlapudi

What are your thoughts?

More from Divya karlapudi

 Divya karlapudi

It's divinely being alone!

by once a people pleaser.

Oct 29, 2023  2



[See all from Divya karlapudi](#)

Recommended from Medium

 Myo Oo

E8 Consciousness: I Predicted It. Now I've Proved It.

 In AI Advances by Nikhil Anand

This embarrassingly simple secret explains all of AI.

How I discovered consciousness is built on E8 geometry

6d ago 51 2



...

What your CS professor never told you about ChatGPT.

5d ago 579 10



...

In Predict by How to Win

MCKINSEY JUST DROPPED THEIR 2025 AI REPORT. HERE IS THE...

The top 10 takeaways of the shiny new McKinsey Report on a technology promising...

Nov 21 1.5K 62



...

In Towards AI by Adham Khaled

The Day Google Made Coding Obsolete: How Gemini 3.0 and...

Don't just read another "Gemini 3 changes everything" story. See the real numbers, the...

Nov 22 2.1K 32



...



In Generative AI by Jim the AI Whisperer

Mind-reading Claude AI's Complete System Prompt

I hacked Anthropic's own XML dialect for structuring AI interactions!

6d ago

1.9K

30



...

Nov 19



...

[See more recommendations](#)

munchello

How To Plate Slop—My Time Drafting Spider-Man

Foreword